

Benchmarking Embedded Databases on IoT Sensor Data

Ethan Rogge, ...

December 2025

1 Introduction to the Technology: Embedded Databases

Embedded databases constitute a class of data management systems designed to run *inside* a host program rather than as an external server process. The DBMS executes within the same address space as the host application and is invoked via direct API calls instead of network protocols. This architectural choice profoundly shapes both the capabilities and limitations of embedded database systems.

1.1 Definition of Embedded Databases

An embedded database is a DBMS that:

- operates as a software library linked to an application or process,
- does not require a separate server or daemon,
- stores data locally on the device's filesystem or in memory,
- handles queries or operations through a direct programming interface.

This contrasts with traditional server-based DBMS—such as PostgreSQL or MySQL—where applications connect via network protocols to an external process that manages concurrency, storage, and query execution.

Embedded databases offer several key advantages:

- **Low latency:** no network round-trip is required.
- **Minimal overhead:** no inter-process communication.
- **Portability:** the database engine is distributed with the binary.
- **Small footprint:** suitable for embedded devices, IoT nodes, or mobile apps.
- **Tight control:** the host program can directly manage caching and storage.

However, these systems also come with inherent limitations:

- reduced concurrency and multi-user capabilities,
- limited or absent SQL support in many engines,
- no access control or advanced security features,
- responsibility for durability and schema evolution often shifted to the host program.

These trade-offs make embedded databases particularly suitable for resource-constrained environments such as IoT nodes, sensors, robotics, or standalone desktop analytics workflows.

1.2 Architectural Models of Embedded Databases

Embedded database technologies vary significantly in their internal design. Several architectural models are commonly used:

1.2.1 Page-oriented storage engines

These systems manage their own page cache, buffer pools, and on-disk B-tree structures. Examples include BerkeleyDB, SQLite, and LMDB. Their properties:

- efficient key-based lookups,
- transactional semantics,
- predictable read/write performance,
- mature caching and recovery mechanisms.

1.2.2 Log-structured merge trees (LSM)

Engines such as LevelDB or RocksDB use LSM trees to support:

- extremely fast writes through append-only logs,
- background compaction processes,
- excellent performance for streaming or time-series ingestion.

These engines are widely used in mobile apps, IoT gateways, and distributed storage systems requiring sustained write throughput.

1.2.3 In-process analytical engines

Modern embedded engines such as DuckDB embed a complete analytical SQL engine directly inside the host process. Their distinguishing features include:

- columnar storage,
- vectorised execution and query pipelining,
- efficient analytical queries (aggregations, joins, window functions),
- no need for a separate database server.

Such systems illustrate the evolution of embedded technology: beyond simple key-value stores, they support full SQL analytics locally within the process.

1.3 Query Capabilities in Embedded Systems

Because embedded engines target environments with constrained resources or specific workloads, their query capabilities vary widely.

Record-oriented APIs. Many engines expose operations such as:

```
get(key),  put(key,value),  delete(key).
```

These APIs offer very fast local access and are ideal for caching or metadata storage.

Full SQL support. Some embedded DBMS such as SQLite or DuckDB support full SQL semantics. This enables local execution of analytical queries, joins, and aggregations without a remote server. Such systems are attractive for standalone desktop analytics and local data processing pipelines.

Hybrid approaches. In practice, systems often combine:

- a lightweight key–value store for fast caching or ingestion,
- a more expressive embedded or external engine for analytical queries.

Our benchmark design reflects this hybrid view by including both a key–value embedded store and an embedded analytical engine, together with a traditional relational baseline.

1.4 Embedded Databases in IoT and Edge Computing

IoT devices frequently require real-time data processing close to where data is generated. Embedded databases are central to these architectures because they allow devices or gateways to:

- store recent sensor readings locally,
- perform preliminary filtering or anomaly detection,
- operate autonomously during network outages,
- reduce bandwidth by transmitting only summarised data,
- respect privacy constraints by keeping raw data on device.

Edge computing architectures often include:

- microcontrollers or single-board computers running a lightweight embedded store,
- intermediate nodes performing local analytics,
- cloud or server components performing long-term storage or heavy computations.

Thus, embedded databases form the backbone of many modern distributed sensing and monitoring systems and provide a natural target for benchmarking in the context of IoT sensor workloads.

1.5 Relevance to Our Benchmark Scenario

The Intel Lab sensor dataset used in our study represents a typical IoT scenario involving time-series measurements collected by multiple devices. From this scenario, we derive a family of benchmark queries that require:

- fast access to recent readings,
- statistical computations per sensor,
- detection of anomalies or outages,
- global analytics over large volumes of historical data.

Our goal is not to build a full application but to design a benchmark workload that captures these requirements realistically and to evaluate how different embedded and non-embedded DBMS behave under this workload.

2 Introduction to the Tools

We now present the three database management systems evaluated in the benchmark: **BerkeleyDB**, **DuckDB**, and **PostgreSQL**. BerkeleyDB and DuckDB are embedded systems, illustrating two different implementations of embedded technology, while PostgreSQL serves as a traditional server-based relational baseline.

Each subsection explains the internal architecture of the tool, its execution model, and its strengths and limitations with respect to sensor data workloads.

2.1 BerkeleyDB

BerkeleyDB (BDB) is a mature, high-performance embedded storage engine that provides key-value access through a simple API. It operates as a library linked directly into the host process, without any external server. As one of the earliest and most widely deployed embedded databases, it has been used in operating systems, networking equipment, web browsers, and mobile devices.

2.1.1 Architecture

BerkeleyDB is built around a **page-oriented storage engine** and uses B-trees as its primary on-disk structure. Data pages are cached in an internal buffer pool, and the engine manages its own journaling, logging, and recovery mechanisms. It exposes a simple interface:

```
put(key, value),  get(key),  delete(key),
```

with optional transactional semantics.

Since the DBMS runs within the same process as the client code, all operations bypass the overhead of inter-process communication. This leads to extremely low latency for key lookups and updates.

2.1.2 Core Features

The main features of BerkeleyDB include:

- **High performance** for point reads, inserts and updates.
- **Configurable storage models**, including B-tree, hash tables, and queues.
- **Transaction support** with ACID guarantees.
- **Crash recovery** through write-ahead logging.
- **Small footprint**, making it suitable for embedded systems and IoT devices.

2.1.3 Transaction Model and Concurrency

BerkeleyDB supports multiversion concurrency control (MVCC) and fine-grained locking. However, concurrency is generally limited to threads within the same process. Unlike server-based DBMS, BerkeleyDB does not natively serve multiple external clients simultaneously.

2.1.4 Limitations

Despite its performance advantages, BerkeleyDB lacks:

- SQL query capabilities,
- built-in analytical operators (aggregations, joins, windows),
- multi-user concurrency architecture,
- ad-hoc querying capabilities.

Any form of analytics must be implemented manually in the client code. For this reason, BerkeleyDB is best viewed as a **local cache or storage component** rather than a general-purpose analytical DBMS.

2.1.5 Relevance to Sensor Workloads

For sensor data, BerkeleyDB is particularly well suited to:

- storing recent readings keyed by sensor identifier and time,
- retrieving the latest values for a given sensor,
- updating or deleting individual measurements,
- maintaining small on-device caches.

Our benchmark uses BerkeleyDB to evaluate how such key-value engines perform on ingestion and point-access phases derived from sensor workloads.

2.2 DuckDB

DuckDB is a modern embedded database designed specifically for analytical workloads. It provides a full SQL interface and implements many of the optimisations found in columnar analytical systems, but packaged as a lightweight in-process library.

2.2.1 Architecture

DuckDB follows a **columnar storage model** and uses vectorised execution. Instead of processing one row at a time, DuckDB processes data in fixed-size vectors (e.g., 1024 rows), greatly improving CPU efficiency and cache locality. It employs late materialisation, compressed storage, and pipelined execution to optimise analytical queries.

The engine runs entirely within the host process, eliminating the need for a database server and reducing overhead for local analytical tasks.

2.2.2 Vectorized Execution Model

Vectorised execution allows DuckDB to:

- evaluate expressions over batches of rows,
- exploit SIMD instructions,
- reduce interpretation overhead,
- pipeline operations efficiently.

This model is particularly advantageous for scans, aggregations, joins, and window functions, making DuckDB well suited for analytical workloads on sensor time-series.

2.2.3 Columnar Storage and Query Processing

DuckDB stores tables in columnar format, meaning that values of the same column are stored contiguously. This enables:

- faster sequential scans for analytical queries,
- data compression,
- better I/O efficiency,
- reduced memory footprint.

Unlike server DBMS such as PostgreSQL, DuckDB does not maintain row-oriented storage optimised for transactional workloads. This explains its poor performance on UPDATE, DELETE, and frequent point-lookups.

2.2.4 Strengths and Limitations

Strengths.

- State-of-the-art analytical performance.
- Fully embedded, no server required.
- Full SQL support including window functions.
- Highly portable and easy to integrate into analysis pipelines.

Limitations.

- Very slow for write-heavy workloads (UPDATE, DELETE).
- Poor performance for random point-lookups.
- Not a multi-user or concurrent transactional system.

2.2.5 Role in the Benchmark

In our benchmark, DuckDB is used as an embedded analytical engine executing aggregation, anomaly detection, outage detection and correlation queries directly on sensor data.

2.3 PostgreSQL (Baseline Reference)

Although PostgreSQL is not an embedded database, it serves as a baseline reference for traditional relational DBMS performance. This allows us to contrast embedded engines with a mature, feature-rich server-based system.

2.3.1 Architecture Overview

PostgreSQL follows a **client–server architecture**. The server manages processes, concurrency control, storage, and query execution. Clients connect through a network protocol, even when running on the same machine.

PostgreSQL stores data in row-oriented format and employs a cost-based optimizer to generate execution plans.

2.3.2 Execution Model and Storage Engine

Key architectural components include:

- multiversion concurrency control (MVCC),
- write-ahead logging (WAL),
- B-tree indexing,
- parallel query execution,
- extensive SQL support.

2.3.3 Strengths and Weaknesses

Strengths.

- Robust ACID transactions.
- Rich indexing options.
- Strong support for SQL standards.
- Good performance for mixed workloads.

Weaknesses.

- Higher overhead due to server architecture.
- Slower analytical scans compared to columnar engines.
- Less suitable for in-process embedded use.

2.3.4 Role as Baseline

PostgreSQL provides a reference point to assess:

- where embedded DBMS outperform server-based architectures,
- where embedded engines fall short,
- how architectural choices influence performance on sensor workloads.

3 Benchmark Design and Experimental Setup

In this section, we describe the dataset, the workload scenario, and the benchmark methodology used to evaluate BerkeleyDB, DuckDB, and PostgreSQL. The focus is exclusively on the benchmark; no full end-user application is implemented. Instead, the benchmark scripts execute a family of queries and updates derived from a realistic IoT monitoring scenario.

3.1 Dataset Description

We use the publicly available Intel Berkeley Research Lab dataset, a widely studied real-world IoT dataset. It contains:

- over 2.3 million timestamped readings,
- collected from 54 sensors deployed indoors,
- covering temperature, humidity, light intensity, and voltage,
- timestamped both as epoch and as human-readable date/time.

Each row corresponds to the measurement of a single sensor at a specific time. This dataset is particularly suited to evaluating embedded databases because:

- it resembles typical IoT monitoring workloads,
- it is large enough to stress test analytical engines,
- it requires both fast point lookups and global computations.

For the benchmark, the raw `data.txt` file is converted to CSV and loaded into the different database engines. To simulate deployments of various scales, subsamples of size 1 000, 10 000, and 100 000 rows are created using a stratified sampling procedure that preserves the chronological order within each sensor.

3.2 Workload Scenario and Requirements

The benchmark workload is inspired by IoT monitoring use cases in which a sensor platform must:

- ingest new measurements,
- retrieve recent values for specific sensors,
- update or delete measurements,
- compute statistics per sensor and per day,
- detect anomalous measurements,
- detect outages (periods with missing data),
- compare the behaviour of different sensors.

These requirements lead to a set of benchmark phases that jointly cover both OLTP-like and OLAP-like operations on the same dataset.

3.3 Benchmark Methodology

To evaluate the performance of the three systems, we designed a benchmark whose phases are directly derived from the workload scenario above. Our objective is to study:

- how each system scales with increasing dataset size,
- whether the behaviour is linear or exponential,
- how embedded systems compare to a traditional RDBMS.

3.3.1 Workload Definition

Each benchmark phase corresponds to a well-defined operation on the sensor dataset. Below we describe in detail what each phase does:

- **LOAD**: bulk insertion of N rows from the CSV file into the database. For BerkeleyDB, each row is stored as a value keyed by a composite key (sensor identifier and timestamp). For DuckDB and PostgreSQL, the rows are inserted into a relational table with columns for timestamp, sensor id, temperature, humidity, light and voltage. This phase models initial ingestion or batch loading.
- **READ**: a series of point lookups of individual rows. The benchmark repeatedly selects specific sensor/timestamp keys and retrieves the corresponding measurement. This models access to individual measurements, as might occur for debugging or detailed inspection of specific events.
- **LAST READINGS**: for a set of sensor identifiers, retrieve the most recent reading of each sensor. In SQL systems, this is implemented as a grouped query using `MAX(timestamp)` or window functions; with BerkeleyDB, it corresponds to accessing the last key for each sensor. This models dashboards or monitoring tools that show the latest state of each sensor.
- **UPDATE**: modify specific fields (e.g. temperature or humidity) of existing rows selected by sensor id and time range. In SQL systems, this is performed with `UPDATE` statements; in BerkeleyDB, the corresponding values are overwritten. This phase simulates corrections to recorded measurements.
- **DELETE**: remove rows corresponding to a subset of sensors or a given time window. In SQL systems, this uses `DELETE` with predicates; in BerkeleyDB, it deletes keys. This models the purging of obsolete or corrupted data.
- **AVG_SENSOR**: compute the average temperature per sensor over all measurements in the dataset. In SQL, this is a `GROUP BY` query: for each sensor id, compute `AVG(temperature)`. This phase quantifies how fast each system can perform basic aggregation over the full dataset.
- **HOT_SENSORS**: identify sensors whose average temperature is above a threshold and that have at least a minimum number of readings. This combines aggregation (`AVG`) with a `HAVING` filter and optionally a `LIMIT k` to obtain a top-k list of the hottest sensors. It models alerting scenarios such as “which sensors are overheating”.
- **ANOMALIES**: detect potentially anomalous measurements for each sensor. Concretely, for each sensor the benchmark computes its mean temperature and either standard deviation or interquartile range, and then selects measurements whose temperature deviates from the mean by more than a fixed threshold (e.g. more than two standard deviations). This is implemented as a combination of aggregations and joins or window functions in SQL systems. It models anomaly detection on individual time-series.
- **DAILY_SUMMARY**: compute daily statistics per sensor, such as the minimum, maximum, and average temperature per day. In SQL, this corresponds to grouping by sensor id and date (extracted from the timestamp) and applying `MIN`, `MAX` and `AVG` aggregations. This models daily reports or summary dashboards.

- **OUTAGES**: detect periods where a sensor did not send any data for an unusually long time. In SQL systems, this is implemented using a window function such as `LAG(timestamp)` partitioned by sensor id and ordered by time; the query computes the difference between consecutive timestamps and filters those where the gap exceeds a given threshold (e.g. one hour). This phase models outage detection in monitoring systems.
- **TOPK_SENSORS**: compute the k sensors with the highest average temperature or with the largest number of readings. This is a `GROUP BY` followed by an `ORDER BY` and `LIMIT`. It models ranking tasks, such as finding the busiest or hottest sensors.
- **COMPARE_TWO**: compare the behaviour of two sensors over time, e.g. by computing the correlation between their temperature time-series. In SQL, this is implemented by joining the measurements of two sensors on aligned timestamps (or time buckets) and then computing correlation coefficients or summary statistics. This models cross-sensor analysis.
- **FULL_PROFILE**: a combined workload that chains several of the analytical queries above. For DuckDB and PostgreSQL, this phase approximates the cost of computing per-sensor averages, daily summaries, outages and top-k sensors in sequence. For BerkeleyDB, which lacks SQL, the benchmark measures a simple local statistical summary over the most recent cached values. This phase illustrates the cumulative cost of a typical analytic session on the dataset.

Together, these phases cover ingestion, point lookups, transactional updates, aggregations, anomaly detection, temporal analysis, and cross-sensor comparisons.

3.3.2 Data Scales

To assess scalability, we evaluate the benchmark using three dataset sizes:

$$N \in \{1\,000, 10\,000, 100\,000\}.$$

For each size, a stratified sampling algorithm selects an approximately equal number of rows per sensor whenever possible and preserves chronological order. This ensures that the benchmark is representative of the full dataset and preserves time-series structure.

3.3.3 Execution Protocol

Each benchmark phase is executed six times consecutively. The first execution is discarded because it typically includes caching, compilation, or buffer pool initialisation effects. The average of the remaining five executions is used as the reported performance metric.

This procedure follows best practices in benchmarking, including:

- warming up caches to avoid cold-start bias,
- averaging over multiple runs to reduce variance,
- ensuring reproducibility across systems.

3.3.4 Hardware and Environment

All experiments were executed on the same machine to guarantee fairness. The benchmark environment includes:

- Python 3.x for orchestrating the benchmark scripts,
- BerkeleyDB with Python bindings,
- DuckDB 1.x,
- a local PostgreSQL server instance,
- identical CSV input files and sampling procedure for all engines.

Details such as CPU model, RAM, operating system, and software versions can be reported to further support reproducibility.

4 Benchmark Results

This section presents the performance obtained for the full set of benchmark phases executed on the three database engines: *BerkeleyDB*, *DuckDB*, and *PostgreSQL*. All experiments were executed for $N \in \{10^3, 10^4, 10^5\}$ rows. Each phase was run six times; the first execution was discarded to avoid cache warm-up effects, and the execution times reported in the figures correspond to the average of the remaining five runs.

We distinguish between (i) OLTP-like operations (LOAD, READ, UPDATE, DELETE, LAST_READINGS) and (ii) analytical SQL operations (AVG_SENSOR, HOT_SENSORS, ANOMALIES, DAILY_SUMMARY, OUTAGES, TOPK_SENSORS, COMPARE_TWO, FULL_PROFILE).

4.1 OLTP-like operations

LOAD. Figure 1 shows the time to load N rows into each DBMS. DuckDB exhibits very poor scalability, increasing from a few seconds at 10^3 rows to more than 330 s for 10^5 rows. This linear but steep growth is due to its immutable storage model and internal checkpointing during bulk inserts. PostgreSQL remains efficient, staying below 1.5 s even for 10^5 rows thanks to optimized WAL handling and bulk insert mechanisms. BerkeleyDB is by far the fastest (below 0.1 s for all N), illustrating the efficiency of a key-value style engine for append-heavy workloads.




Figure 1: LOAD — average duration vs. number of rows.

READ. The READ benchmark retrieves records by key (sensor identifier and, where applicable, timestamp). As shown in Figure 2, BerkeleyDB dominates with sub-millisecond response times. DuckDB and PostgreSQL both scale linearly, but DuckDB becomes extremely slow at 10^5 rows (more than 70 s), confirming that it is not optimized for repeated point lookups. PostgreSQL remains significantly faster (below 26 s), relying on index-based lookups.




Figure 2: READ — average duration vs. number of rows.

UPDATE and DELETE. Figures 3 and 4 show that BerkeleyDB again provides the best performance with almost constant-time updates and deletions. PostgreSQL is slower but scales predictably, reaching around 17 s for UPDATE and 3 s for DELETE at 10^5 rows. DuckDB is by far the slowest (over 240 s for UPDATE and 11 s for DELETE at 10^5), due to table rewrites induced by its append-only columnar storage.




Figure 3: UPDATE — average duration vs. number of rows.




Figure 4: DELETE — average duration vs. number of rows.

LAST_READINGS. This phase retrieves the most recent readings for a set of sensors. Figure 5 shows that BerkeleyDB is the fastest (fractions of milliseconds even at 10^5 rows). DuckDB performs moderately well (≈ 21 ms at 10^5), while PostgreSQL is the slowest (≈ 31 ms at 10^5), though still within acceptable limits.




Figure 5: LAST_READINGS — average duration vs. number of rows.

4.2 Analytical SQL operations

ANOMALIES. The ANOMALIES phase computes deviations from the mean temperature for each sensor using SQL aggregations and window or join logic. As shown in Figure 6, DuckDB is consistently faster (from 0.006 s to 0.022 s) than PostgreSQL (from 0.002 s to 0.070 s), thanks to vectorised columnar execution. BerkeleyDB cannot execute this query directly.




Figure 6: ANOMALIES — performance comparison between DuckDB and PostgreSQL.

AVG_SENSOR. Figure 7 presents the average execution time to compute the average temperature per sensor via a GROUP BY query. DuckDB clearly outperforms PostgreSQL, achieving

about 0.01 s at 10^5 rows, versus 0.027 s for PostgreSQL.




Figure 7: AVG_SENSOR — performance comparison.

HOT_SENSORS. The HOT_SENSORS query returns sensors whose average temperature exceeds a threshold and have enough readings. As seen in Figure 8, both systems scale linearly, but DuckDB is systematically faster than PostgreSQL.




Figure 8: HOT_SENSORS — performance comparison.

OUTAGES. The OUTAGES phase detects large gaps in the time series using a LAG window function on timestamps. Figure 9 shows that DuckDB scales from roughly 0.003 s to 0.014 s, while PostgreSQL increases from 0.001 s to about 0.040 s. Both are linear, but DuckDB has a smaller slope.




Figure 9: OUTAGES — performance comparison.

TOPK_SENSORS. For the TOPK_SENSORS query (top sensors by average temperature), Figure 10 confirms DuckDB’s advantage: its execution time remains around three times lower than PostgreSQL, while both curves grow linearly with N .




Figure 10: TOPK_SENSORS — performance comparison.

COMPARE_TWO. The COMPARE_TWO query computes a correlation-like measure between the temperature time-series of two sensors by joining their measurements on aligned timestamps. Figure 11 shows that DuckDB remains fast (≈ 0.005 s at 10^5), whereas PostgreSQL is considerably slower (≈ 0.029 s).




Figure 11: COMPARE_TWO — performance comparison.

DAILY_SUMMARY. Figure 12 shows that PostgreSQL is faster than DuckDB for the DAILY_SUMMARY aggregation. DuckDB reaches about 0.050 s at 10^5 rows, while PostgreSQL stays near 0.023 s. This can be explained by differences in how both engines materialise intermediate grouped results.




Figure 12: DAILY_SUMMARY — performance comparison.

FULL_PROFILE. Finally, Figure 13 presents the FULL_PROFILE workload. For DuckDB and PostgreSQL, this corresponds to the combined cost of several analytical queries (AVG_SENSOR, DAILY_SUMMARY, OUTAGES, TOPK_SENSORS), whereas BerkeleyDB executes only a local statistical summary on its cached data. DuckDB provides the best analytical performance (0.012 s to 0.078 s), PostgreSQL is slower (0.018 s to 0.105 s), and BerkeleyDB is extremely fast for its limited local task (0.004 s to 0.030 s).




Figure 13: FULL_PROFILE — combined profile workload.

5 Discussion

The experimental results highlight a clear separation of responsibilities between the three engines tested. Each system exhibits strengths that are consistent with its architectural design.

5.1 BerkeleyDB

Strengths. BerkeleyDB offers outstanding performance for key–value access patterns. It provides near-constant-time READ, UPDATE, and DELETE operations, and extremely fast LOAD times even for 10^5 rows. Its footprint is small and it can be embedded directly in host programs, making it an excellent candidate for resource-constrained IoT devices or gateways.

Weaknesses. BerkeleyDB does not support SQL and cannot execute analytical queries such as aggregations, joins, or window functions. All higher-level logic must be implemented in the client code. As a result, BerkeleyDB is unsuitable as a standalone analytical database and must be complemented by another engine for global analysis.

5.2 DuckDB

Strengths. DuckDB delivers the best performance for all analytical queries in our benchmark (ANOMALIES, AVG_SENSOR, HOT_SENSORS, OUTAGES, TOPK_SENSORS, COMPARE_TWO, and the analytical part of FULL_PROFILE). Its columnar storage and vectorised execution make it particularly well adapted to OLAP workloads on large sensor datasets, while still running as an embedded library without a separate server process.

Weaknesses. On the other hand, DuckDB performs extremely poorly on OLTP-like workloads: LOAD, UPDATE, DELETE and repeated point READs are orders of magnitude slower than with BerkeleyDB or even PostgreSQL. This is a direct consequence of its append-only storage and optimisation for scans rather than random writes. It is therefore not suitable for write-heavy or latency-sensitive transactional workloads.

5.3 PostgreSQL

Strengths. PostgreSQL provides balanced behaviour across workloads. It is much faster than DuckDB for UPDATE and DELETE operations, and its index-based execution plans offer reasonable performance for READ and LAST_READINGS. For analytics, it is systematically slower than DuckDB, but still scales linearly and remains within low tens of milliseconds for 10^5

rows. As a mature relational system, it also offers full transactional guarantees, rich indexing options and a robust client–server architecture.

Weaknesses. PostgreSQL is not optimised for columnar scans and cannot match DuckDB on pure analytical workloads. It also has a higher operational overhead (separate server process, configuration, memory usage) than embedded engines such as DuckDB or BerkeleyDB.

5.4 Overall assessment

Taken together, the results show that no single engine dominates all workloads. Instead, each engine is optimal for a different role in sensor data processing:

- **Device / edge role:** BerkeleyDB is ideal as a local cache of recent readings, thanks to its extremely fast key–value access and tiny footprint.
- **Embedded analytics role:** DuckDB is the best choice for complex analytical queries and batch processing over large sensor datasets.
- **Server / baseline role:** PostgreSQL provides a robust, general-purpose relational backend and serves as a solid reference for comparison.

The benchmark confirms that embedded key–value technology is extremely well suited for local, latency-critical access, but must be complemented by an analytical engine to support the full range of sensor analytics workloads.