

# Benchmarking Embedded Databases on IoT Sensor Data

INFO-H-415 Advanced Databases

DON Basem  
CHEOURIFA Anas  
ROGGE Ethan

ZIMANYI Esteban

# Outline

- Motivation of this benchmark
- Berkeley DB low-level, key-value embedded engine
- DuckDB analytical, in-process SQL DB
- PostgreSQL used as a baseline
- Methodology
- Key results ( OLTP / OLAP )
- Comparative takeaways & recommandations

# What is an Embedded Database ?

## Local

No server process, no network communication

## Storage

Data stored locally (files or memory)



## Library

As a library inside the application

## Execution

Direct in-process execution

# Embedded Database

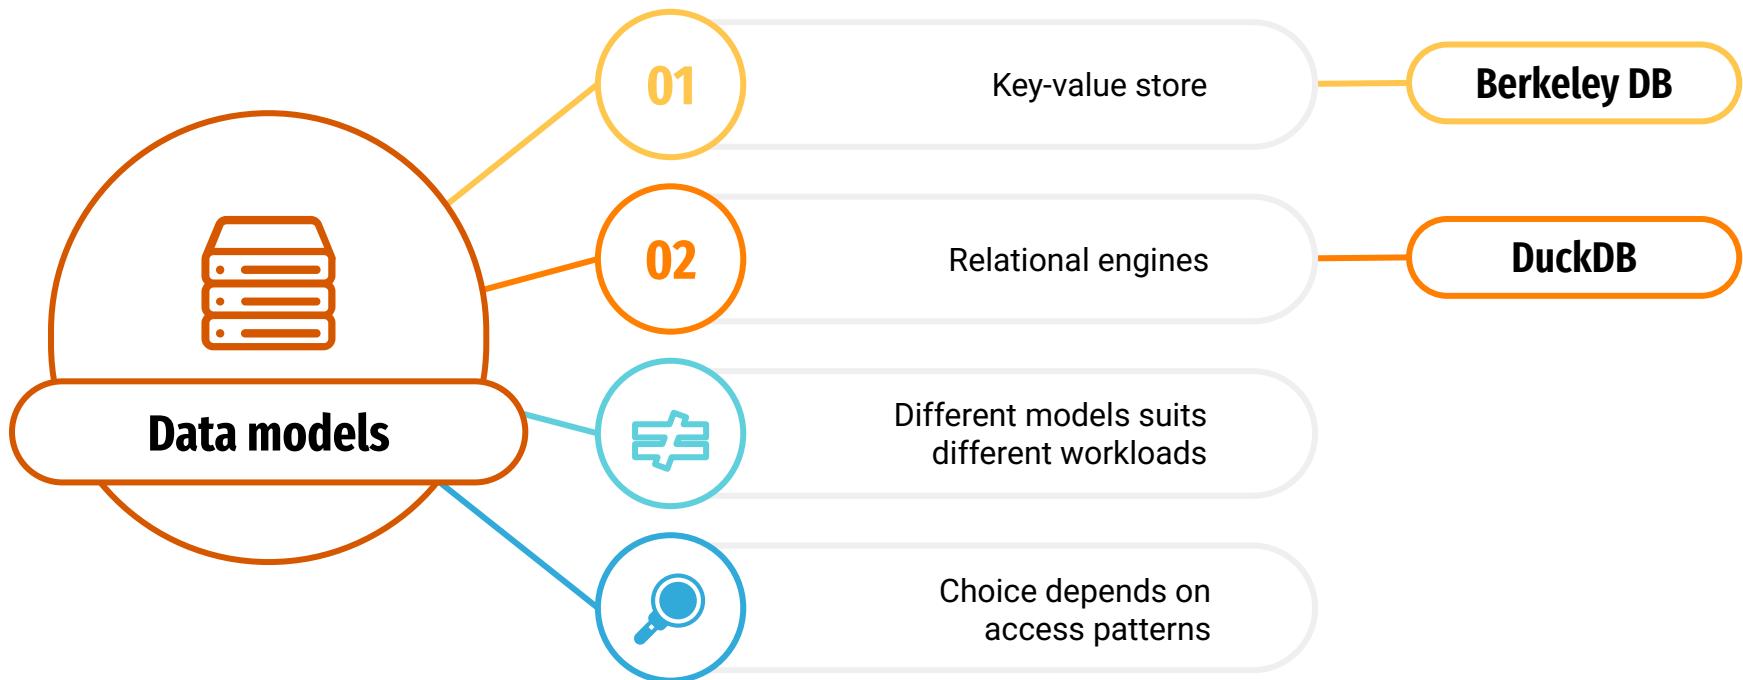
## Advantages

1. Low latency
2. Simple deployment
3. Small footprint

## Trade-offs

1. Limited concurrency
2. Reduced expressiveness
3. Application-level logic

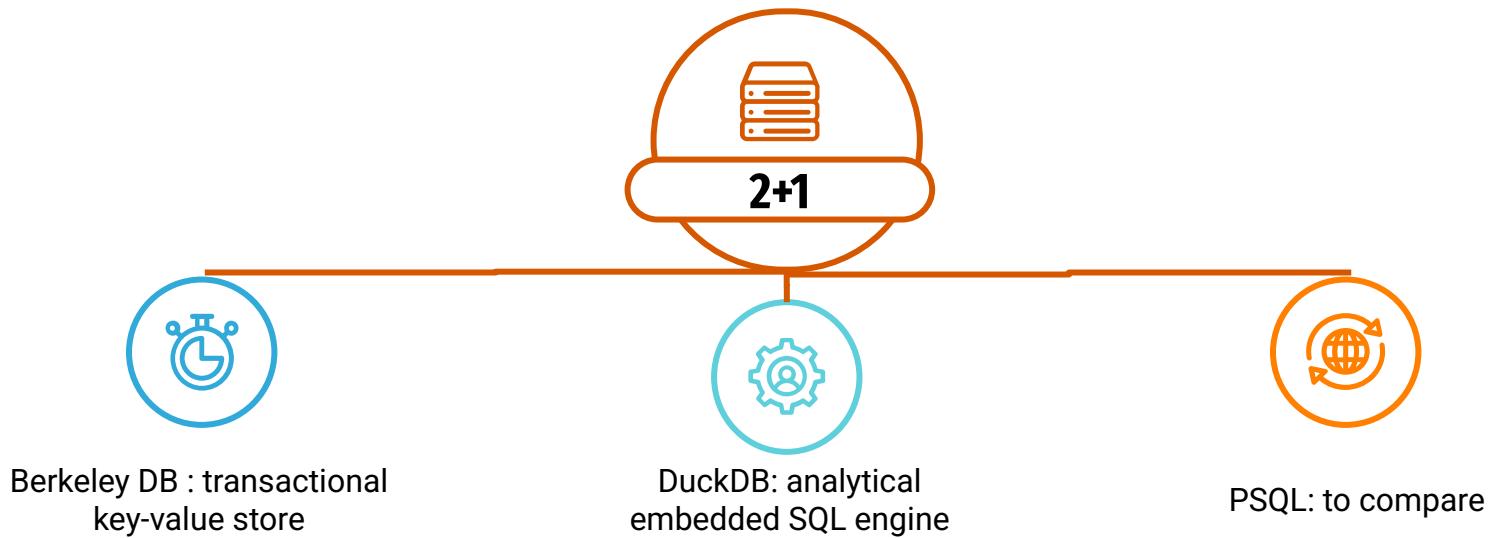
# Data models



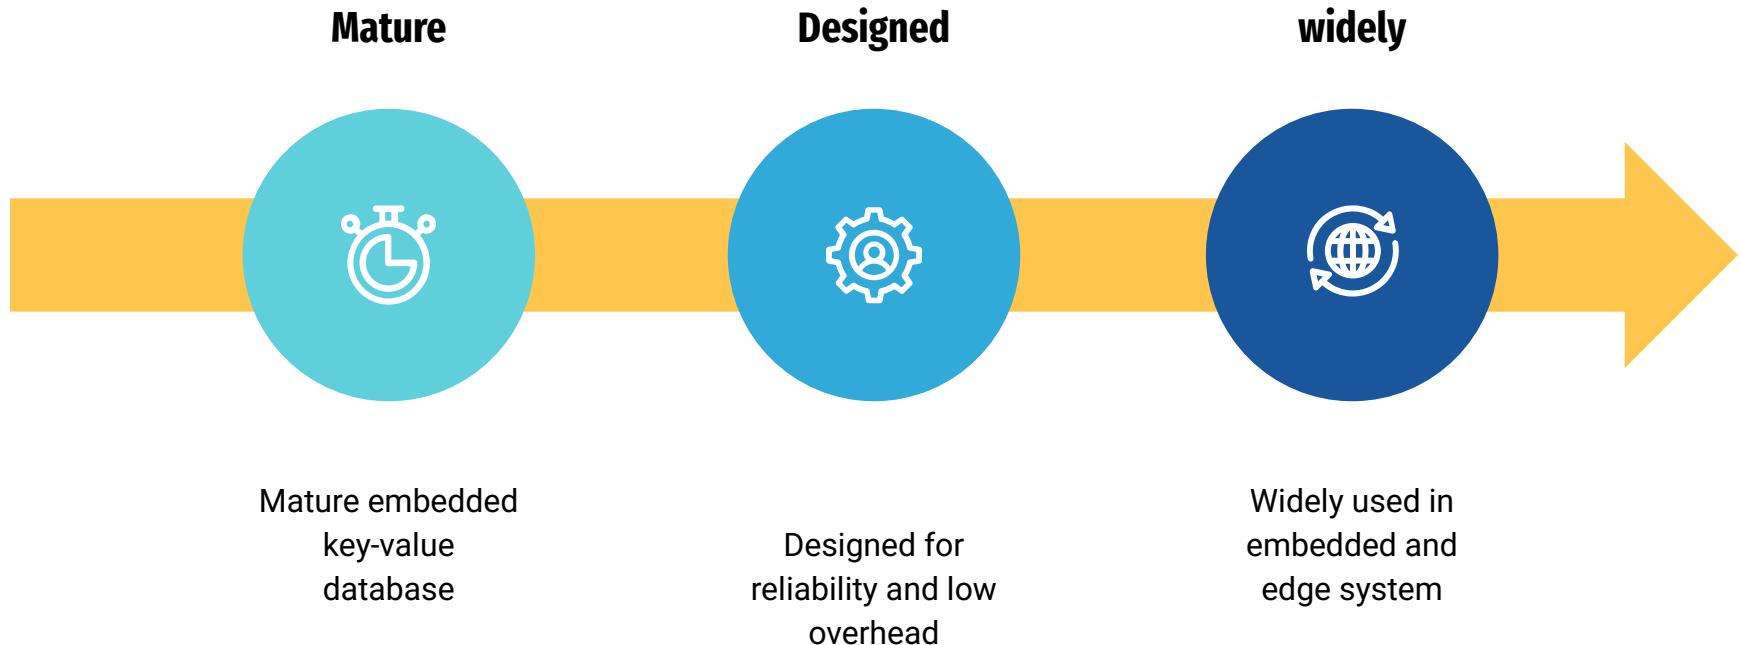
# Why IoT Sensor Data ?



# Tools overview



# Berkeley DB - Positioning



# Berkeley DB - Architecture

B-tree, hash and record-based storage

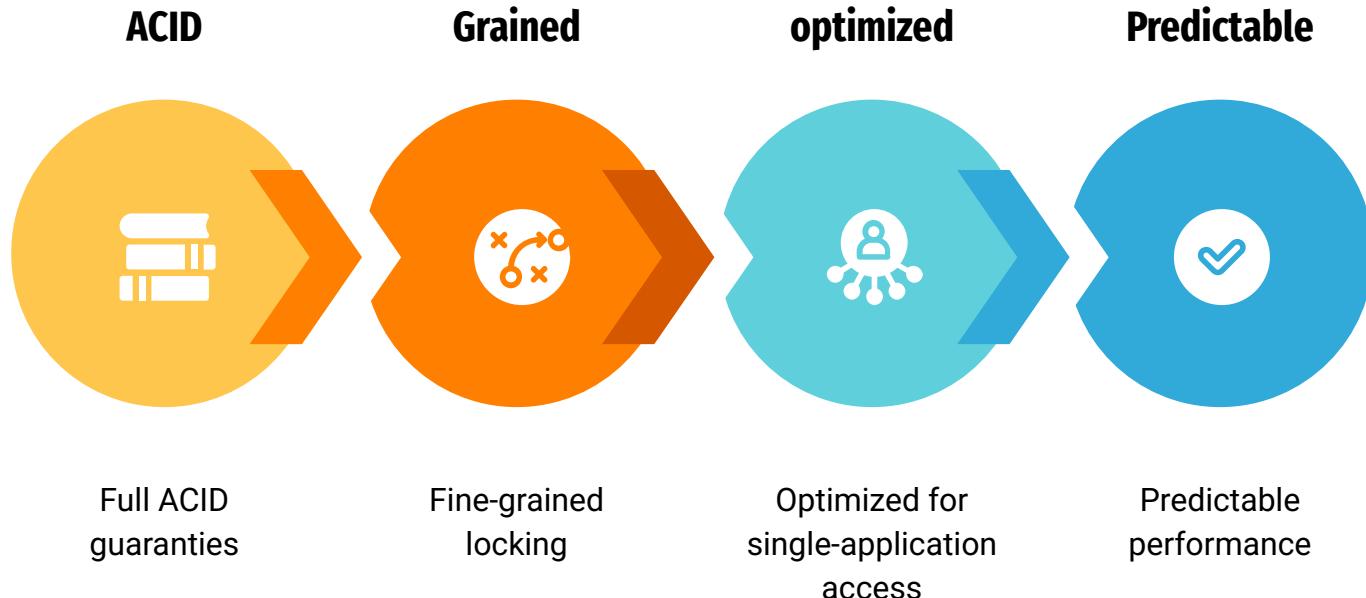
Key-value data model

Write-ahead log (WAL)

Page-based engine with buffer cache



# Berkeley DB - Transactions & Concurrency



# Berkeley DB - Capabilities & Limits

## Excellent

1. Fast insert
2. Point lookups

## Limitations

1. No SQL
2. No build-in analytics

# DuckDB - positioning

## Optimised for OLAP

fast scans, joins, aggregations  
and native Parquet/CSV access

## Rich language bindings

Python, R, Java, Node.js, etc



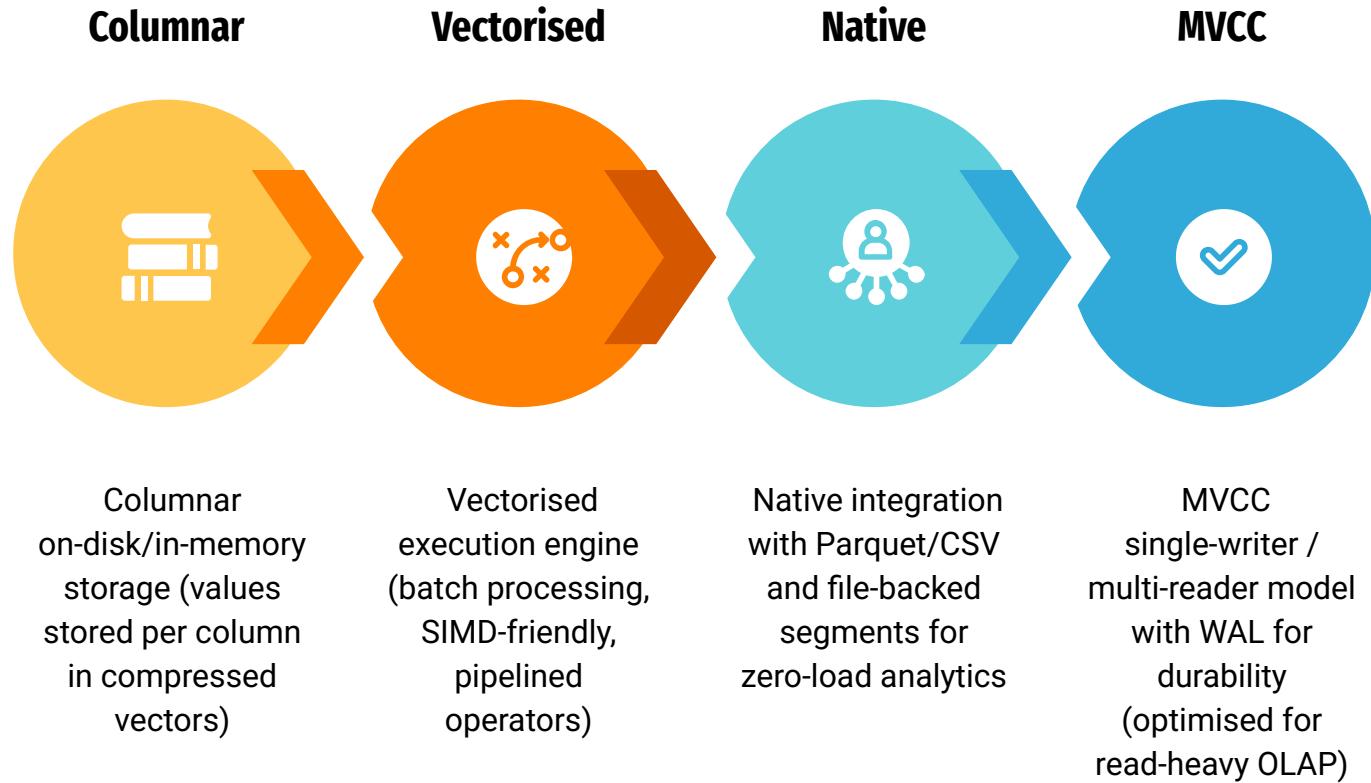
## Modern in-process analytical SQL engine

columnar storage & vectorised  
execution

## Analytical

Ideal for embedded analytics, data  
science and notebook workflows

# DuckDB - Architecture



# DuckDB - Capabilities & Limits

## Excellent for

1. In-process analytical SQL: complex joins, aggregations, window functions
2. Large sequential scans and columnar analytics (fast compression + SIMD)

## Limitations

1. Not designed for high-frequency point updates, deletes or many small random writes
2. Single-writer model (MVCC) limits concurrent writers
3. Higher memory/CPU during heavy analytical queries; less suitable for extremely constrained devices
4. Slower for repeated low-latency point lookups compared to key-value engines

# Dataset

**Intel Berkeley Research  
Lab dataset**

1. Real-world sensor data
2. Temperature, humidity,  
light, voltage
3. Scales: 1K, 10k, 100K,  
1M rows

# Benchmark Methodology

Same dataset and queries

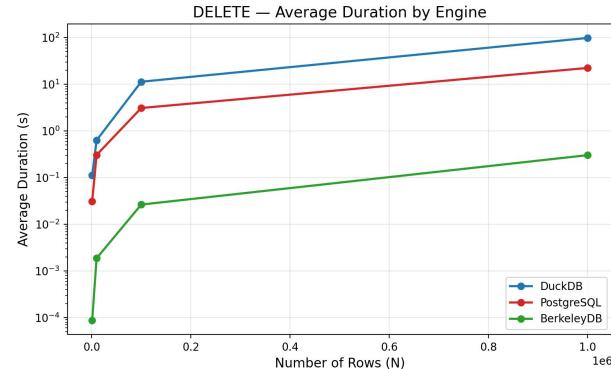
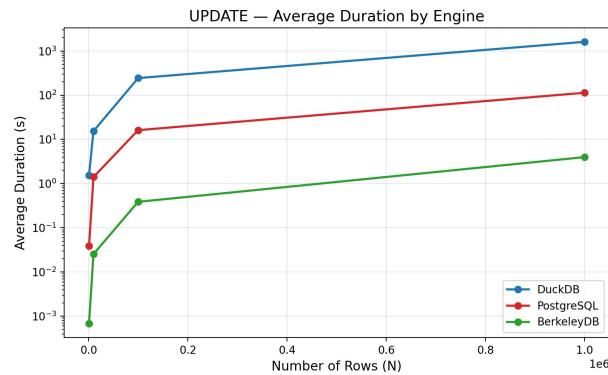
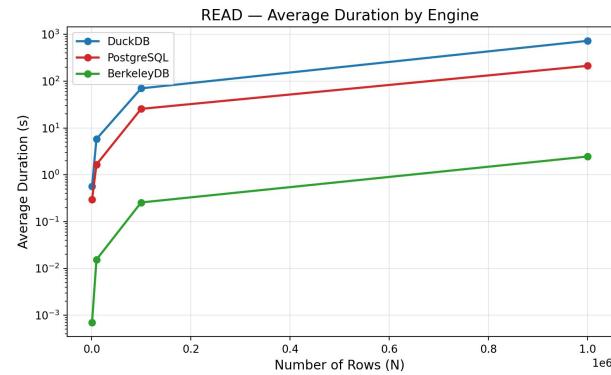
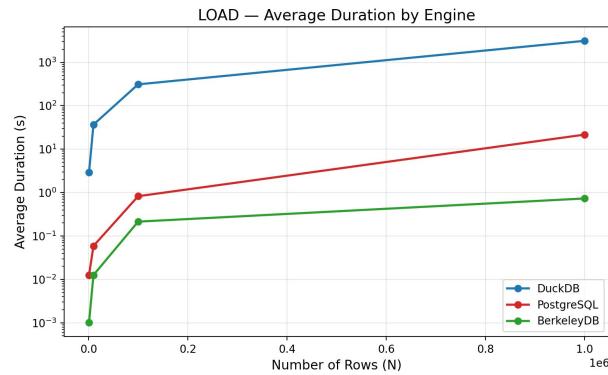
Same machine for all systems

First run discarded, average reported

Each query executed 6 times



# OLTP Queries Results



# OLTP Results - Explanation

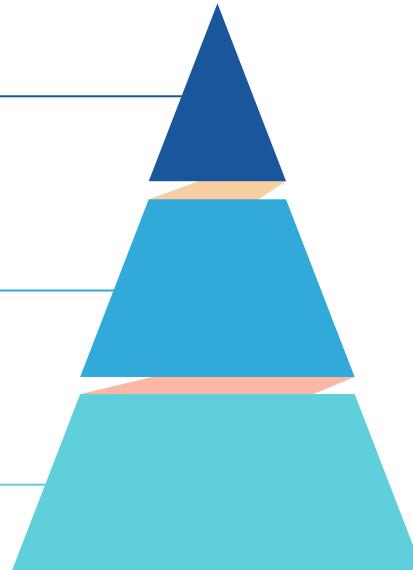
Direct key-value access,  
minimal execution path



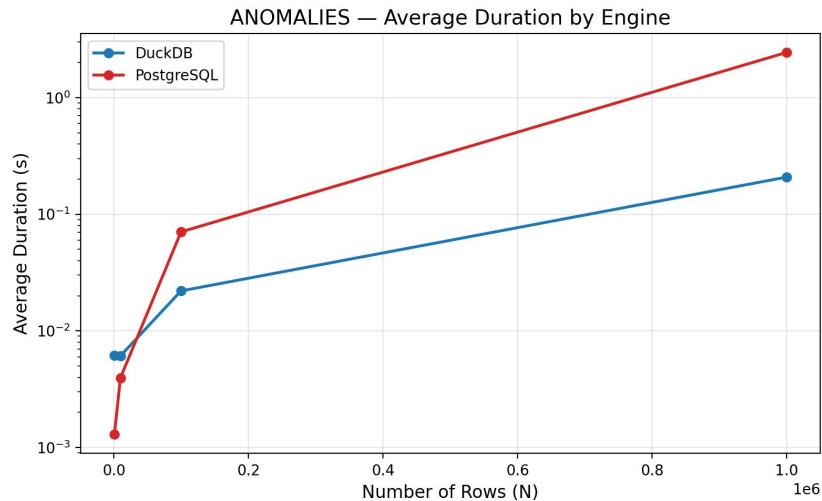
Index-based access but SQL +  
MVCC overhead



Columnar layout -> write  
amplification & poor random  
access

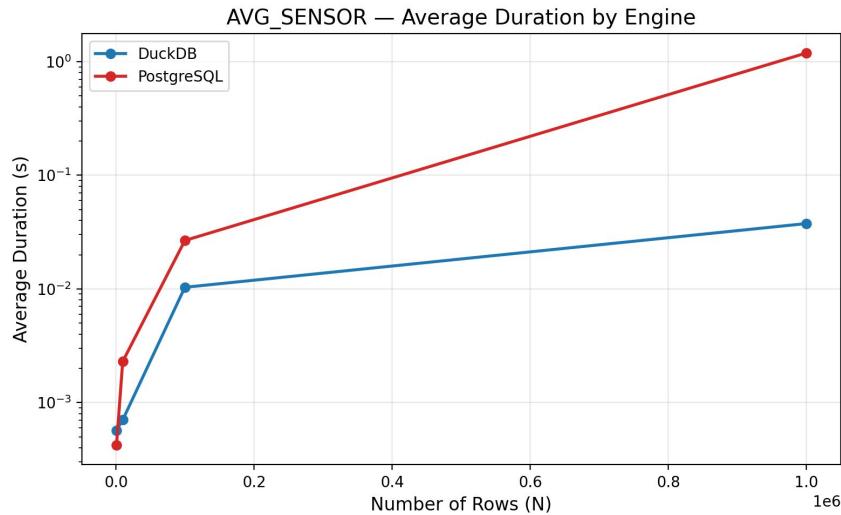


# ANOMALIES



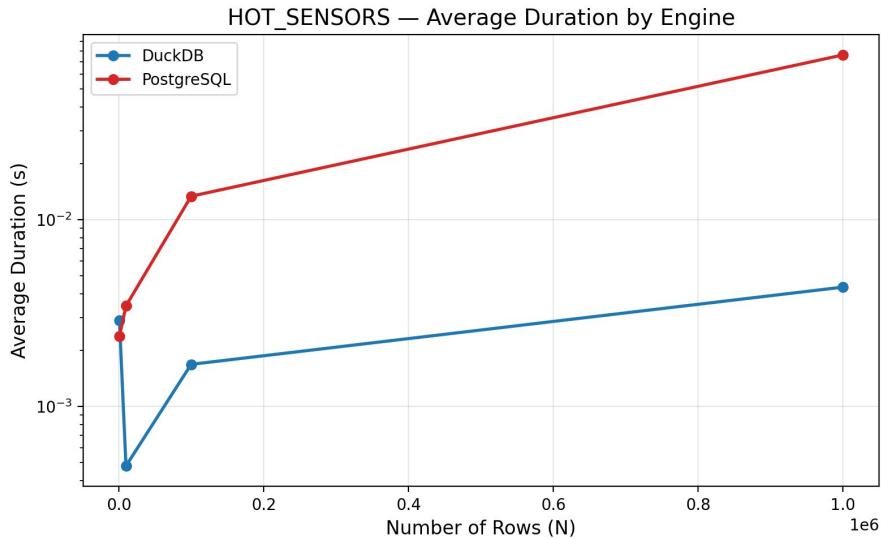
- Combinaison of aggregations and joins or window functions

# AVG\_SENSOR



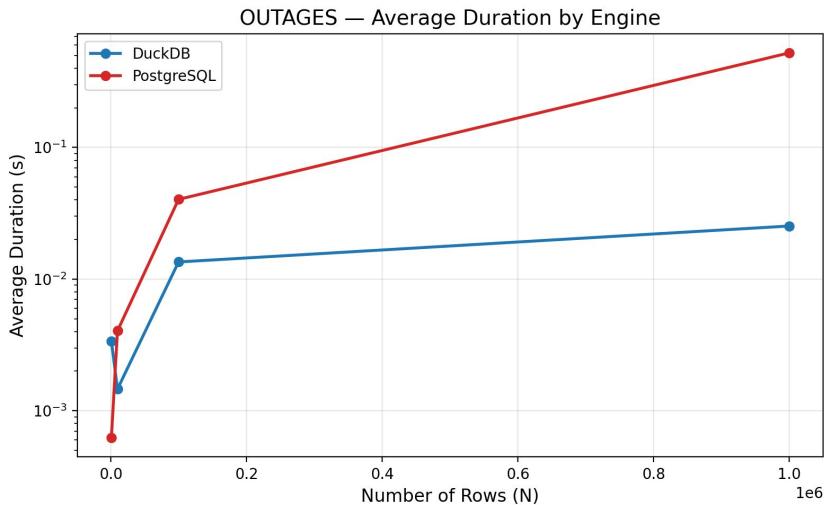
- GROUP BY query: for each sensor id, compute  $\text{AVG}(\text{temperature})$

# HOT\_SENSORS



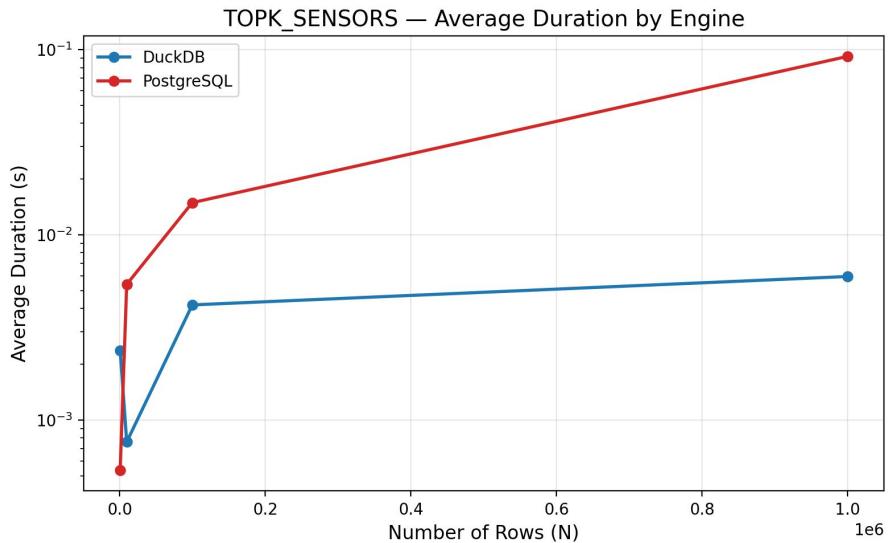
- Combine aggregation (AVG) with a HAVING filter and optionally a LIMIT k

# OUTAGES



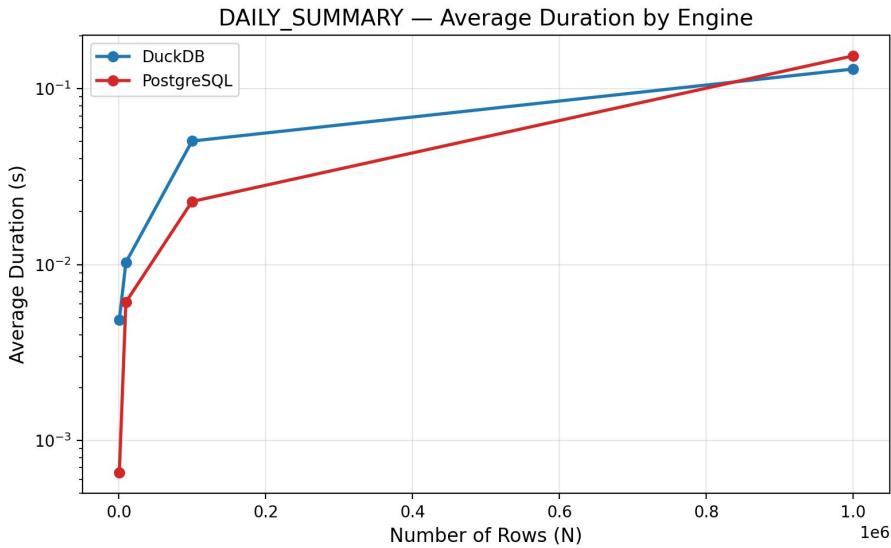
- Window function such as  $\text{LAG(timestamp)}$  partitioned by sensor id and ordered by time

# TOPK\_SENSORS



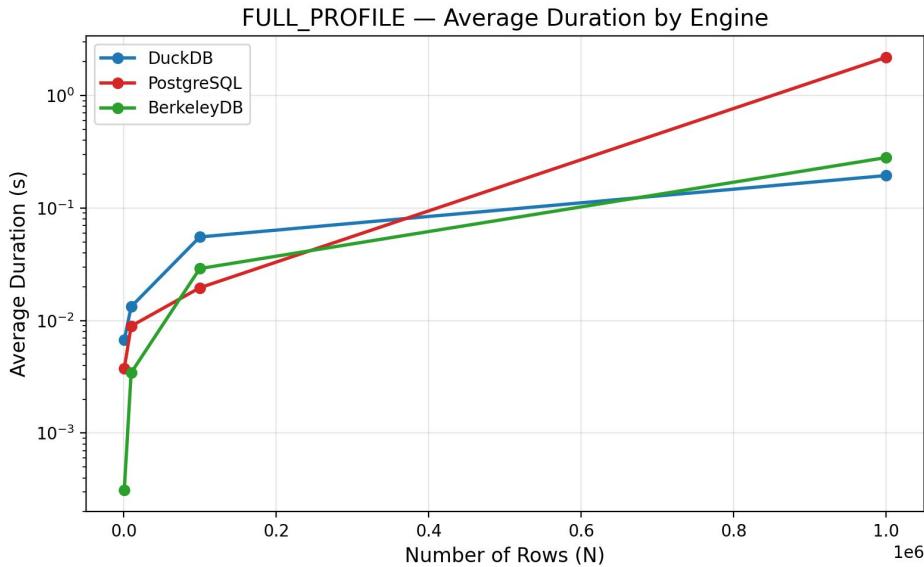
- GROUP BY followed by an ORDER BY and LIMIT

# DAILY\_SUMMARY



- GROUP BY sensor id and date and applying MIN, MAX and AVG aggregations

# FULL\_PROFILE



- DuckDB & PSQL: sum of averages, daily summary, outages and top-k sensors
- Berkeley DB: simple local statistical summary over the most recent cached values

# Analytical SQL Operations Results - Explanation

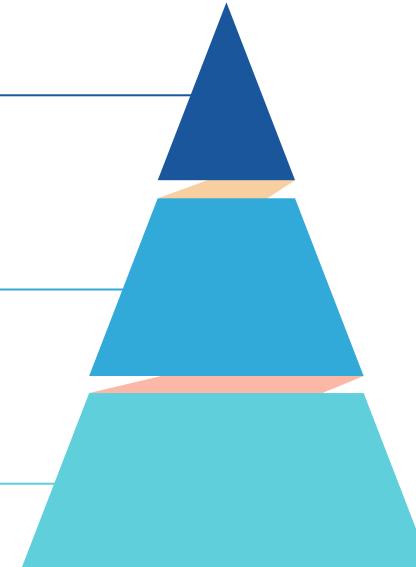
Columnar storage + vectorized execution (SIMD, cache-friendly)



Row-oriented, tuple-at-a-time execution



No native analytical operators



# Comparative Assessment



## Berkeley DB

Best for edge ingestion



## DuckDB

Best for embedded analytics



## PSQL

General-purpose baselinet

# Conclusion

- No single best database
- Architecture determines performance
- Hybrid architectures are often optimal